

q2

December 8, 2020

```
[50]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics
```

0.1 a) load/merge data and visualize logerror

```
[51]: # load data into DataFrames
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

train_df = pd.read_csv("train.csv", keep_default_na=False, na_values=[""])
properties_df = pd.read_csv("properties.csv", keep_default_na=False,
    ↪na_values=[""])


df = pd.merge(train_df, properties_df, on='id')
# [31725 rows x 60 columns]
#X = dataset.iloc[:, [3, 4]].values
```

```
[52]: # eliminate outliers
min_value, max_value = np.percentile(df['logerror'], [1, 99])

df['logerror'] = np.where(df['logerror']>max_value, max_value, df['logerror'])
df['logerror'] = np.where(df['logerror']<min_value, min_value, df['logerror'])
```


```
[53]: # scatter of logerr
fig, ax = plt.subplots()
ax.scatter(df['logerror'], df['logerror'].index)
# plt.show()
```

[53]: <matplotlib.collections.PathCollection at 0x7fb40e1aad0>



```
[54]: # histogram of logerr  
# # df.logerror.hist()  
df.logerror.hist(bins=30)
```

[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb40e0b20d0>



0.2 b) data cleaning

```
[55]: # build new data frame
num = len(df)
data = pd.DataFrame({'column_name': list(df.columns), 'missing_count': list(df.
    ↪isnull().sum())})
data['missing_ratio'] = data.iloc[:,1].values/num
print(data)
```

	column_name	missing_count	missing_ratio
0	id	0	0.000000
1	logerror	0	0.000000
2	transactiondate	0	0.000000
3	airconditioningtypeid	21563	0.679685
4	architecturalstyletypeid	31628	0.996942
5	basementsqft	31711	0.999559
6	bathroomcnt	0	0.000000
7	bedroomcnt	0	0.000000
8	buildingclasstypid	31717	0.999748
9	buildingqualitytypeid	11488	0.362112
10	calculatedbathnbr	414	0.013050
11	decktypeid	31502	0.992971
12	finishedfloor1squarefeet	29381	0.926115
13	calculatedfinishedsquarefeet	228	0.007187
14	finishedsquarefeet12	1647	0.051915
15	finishedsquarefeet13	31711	0.999559
16	finishedsquarefeet15	30454	0.959937
17	finishedsquarefeet50	29381	0.926115
18	finishedsquarefeet6	31591	0.995776
19	fips	0	0.000000
20	fireplacecnt	28374	0.894374
21	fullbathcnt	414	0.013050
22	garagecarcnt	21280	0.670764
23	garagetotalsqft	21280	0.670764
24	hashottuborspa	30929	0.974909
25	heatingorsystemtypeid	11962	0.377053
26	latitude	0	0.000000
27	longitude	0	0.000000
28	lotsizesquarefeet	3522	0.111017
29	poolcnt	25454	0.802333
30	poolsizeum	31394	0.989567
31	pooltypeid10	31337	0.987770
32	pooltypeid2	31317	0.987139
33	pooltypeid7	25862	0.815193

34	propertycountylandusecode	0	0.000000
35	propertylandusetypeid	0	0.000000
36	propertyzoningdesc	11135	0.350985
37	rawcensustractandblock	0	0.000000
38	regionidcity	666	0.020993
39	regionidcounty	0	0.000000
40	regionidneighborhood	19082	0.601481
41	regionidzip	12	0.000378
42	roomcnt	0	0.000000
43	storytypeid	31711	0.999559
44	threequarterbathnbr	27471	0.865910
45	typeconstructiontypeid	31613	0.996470
46	unitcnt	11127	0.350733
47	yardbuildingsqft17	30814	0.971284
48	yardbuildingsqft26	31691	0.998928
49	yearbuilt	260	0.008195
50	numberofstories	24526	0.773081
51	fireplaceflag	31631	0.997037
52	structuretaxvaluedollarcnt	128	0.004035
53	taxvaluedollarcnt	1	0.000032
54	assessmentyear	0	0.000000
55	landtaxvaluedollarcnt	1	0.000032
56	taxamount	1	0.000032
57	taxdelinquencyflag	31112	0.980678
58	taxdelinquencyyear	31112	0.980678
59	censustractandblock	208	0.006556

```
[56]: # fill missing data
df.update(df.fillna(df.mean(), inplace=True))
print(df)
```

	id	logerror	transactiondate	airconditioningtypeid	\
0	14366692	-0.1684	1/1/16	1.809683	
1	14739064	-0.0030	1/2/16	1.809683	
2	10854446	0.3825	1/3/16	1.809683	
3	11672170	-0.0161	1/3/16	1.000000	
4	12524288	-0.0419	1/3/16	1.809683	
...	
31720	12756771	0.0658	12/30/16	1.809683	
31721	11295458	-0.0294	12/30/16	1.000000	
31722	11308315	0.0070	12/30/16	1.000000	
31723	11703478	0.0431	12/30/16	1.809683	
31724	12566293	0.4207	12/30/16	1.809683	

	architecturalstyletypeid	basementsqft	bathroomcnt	bedroomcnt	\
0	7.453608	670.571429	3.5	4.0	
1	7.453608	670.571429	1.0	2.0	
2	7.453608	670.571429	2.0	2.0	

3	7.453608	670.571429	4.0	5.0
4	7.453608	670.571429	1.0	1.0
...
31720	7.453608	670.571429	1.0	3.0
31721	7.453608	670.571429	2.0	2.0
31722	7.453608	670.571429	3.0	5.0
31723	7.453608	670.571429	1.0	3.0
31724	7.453608	670.571429	1.0	3.0

	buildingclasstypeid	buildingqualitytypeid	...	numberofstories	\
0	4.0	5.570193	...	1.450479	
1	4.0	5.570193	...	1.450479	
2	4.0	7.000000	...	1.450479	
3	4.0	1.000000	...	1.450479	
4	4.0	7.000000	...	1.450479	
...	
31720	4.0	7.000000	...	1.450479	
31721	4.0	7.000000	...	1.450479	
31722	4.0	4.000000	...	1.450479	
31723	4.0	7.000000	...	1.450479	
31724	4.0	7.000000	...	1.450479	

	fireplaceflag	structuretaxvaluedollarcnt	taxvaluedollarcnt	\
0	1	346458.0	585529.0	
1	1	66834.0	210064.0	
2	1	55396.0	105954.0	
3	1	559040.0	1090127.0	
4	1	56233.0	70316.0	
...	
31720	1	65728.0	307167.0	
31721	1	40163.0	50203.0	
31722	1	248378.0	331525.0	
31723	1	17520.0	39934.0	
31724	1	66258.0	163037.0	

	assessmentyear	landtaxvaluedollarcnt	taxamount	taxdelinquencyflag	\
0	2015	239071.0	10153.02	NaN	
1	2015	143230.0	2172.88	NaN	
2	2015	50558.0	1443.69	NaN	
3	2015	531087.0	13428.94	NaN	
4	2015	14083.0	913.17	NaN	
...	
31720	2015	241439.0	4038.70	NaN	
31721	2015	10040.0	1263.39	Y	
31722	2015	83147.0	6461.79	NaN	
31723	2015	22414.0	627.91	NaN	
31724	2015	96779.0	2560.96	NaN	

	taxdelinquencyyear	censustractandblock
0	13.314845	6.048996e+13
1	13.314845	6.059040e+13
2	13.314845	6.037140e+13
3	13.314845	6.037260e+13
4	13.314845	6.037570e+13
...
31720	13.314845	6.037550e+13
31721	15.000000	6.037900e+13
31722	13.314845	6.037900e+13
31723	13.314845	6.037230e+13
31724	13.314845	6.037540e+13

[31725 rows x 60 columns]

0.3 c) univariate analysis


```
[57]: # make bar chart
correlation = df.corr()['logerror']
neg = correlation.index.isin(['logerror'])
correlation = correlation[~neg]
print(correlation)
correlation.plot.bar()
```

id	0.006562
airconditioningtypeid	0.006328
architecturalstyletypeid	-0.001234
basementsqft	0.005239
bathroomcnt	0.033445
bedroomcnt	0.032168
buildingclasstypid	NaN
buildingqualitytypeid	-0.001840
calculatedbathnbr	0.034345
decktypeid	NaN
finishedfloor1squarefeet	0.000807
calculatedfinishedsquarefeet	0.042841
finishedsquarefeet12	0.039504
finishedsquarefeet13	0.012608
finishedsquarefeet15	0.014687
finishedsquarefeet50	0.000621
finishedsquarefeet6	-0.000656
fips	0.007863
fireplacecnt	0.005099
fullbathcnt	0.032986
garagecarcnt	-0.000039
garagetotalsqft	0.005227
heatingorsystemtypeid	-0.019511

latitude	0.003277
longitude	0.007782
lotsizesquarefeet	0.006093
poolcnt	NaN
poolsizeum	-0.001442
pooltypeid10	NaN
pooltypeid2	NaN
pooltypeid7	NaN
propertylandusetypeid	-0.028459
rawcensustractandblock	0.007821
regionidcity	-0.000542
regionidcounty	-0.004874
regionidneighborhood	-0.004454
regionidzip	-0.001253
roomcnt	0.009762
storytypeid	NaN
threequarterbathnbr	-0.001133
typeconstructiontypeid	-0.002361
unitcnt	0.005964
yardbuildingsqft17	-0.006088
yardbuildingsqft26	0.004131
yearbuilt	0.016442
numberofstories	0.011565
structuretaxvaluedollarcnt	0.005950
taxvaluedollarcnt	-0.002060
assessmentyear	NaN
landtaxvaluedollarcnt	-0.006758
taxamount	-0.024282
taxdelinquencyyear	-0.011826
censustractandblock	0.008389

Name: logerror, dtype: float64

[57]: <matplotlib.axes._subplots.AxesSubplot at 0x7fb410869280>



1 explain reason

This bar plot above indicates that there is generally either a positive correlation and negative correlation between the logerror and each variable. The few variables without correlation values indicate that the association between those variables and logerror is not obvious or is hardly observable. The variables with positive correlation values indicate that large values of logerror corresponds to large values of those variables, and vice versa. While the variables with negative correlation values indicate that large values of logerror corresponds to small values of those variables, and vice versa.

1.1 d) non-linear regression model

```
[58]: # drop categorical features
# ("hashottuborspa", "propertycountylandusecode", "propertyzoningdesc",
↪ "fireplaceflag", "taxdelinquencyflag")
# drop "id" and "transactiondate"
df.drop(['id', 'transactiondate'], axis=1, inplace=True)
df.fillna(df.mean(), inplace=True)
print(df)
```

	logerror	airconditioningtypeid	architecturalstyletypeid \
0	-0.1684	1.809683	7.453608
1	-0.0030	1.809683	7.453608
2	0.3825	1.809683	7.453608
3	-0.0161	1.000000	7.453608
4	-0.0419	1.809683	7.453608
...
31720	0.0658	1.809683	7.453608
31721	-0.0294	1.000000	7.453608
31722	0.0070	1.000000	7.453608
31723	0.0431	1.809683	7.453608
31724	0.4207	1.809683	7.453608

	basementsqft	bathroomcnt	bedroomcnt	buildingclasstypeid \
0	670.571429	3.5	4.0	4.0
1	670.571429	1.0	2.0	4.0
2	670.571429	2.0	2.0	4.0
3	670.571429	4.0	5.0	4.0
4	670.571429	1.0	1.0	4.0
...
31720	670.571429	1.0	3.0	4.0
31721	670.571429	2.0	2.0	4.0
31722	670.571429	3.0	5.0	4.0
31723	670.571429	1.0	3.0	4.0
31724	670.571429	1.0	3.0	4.0

	buildingqualitytypeid	calculatedbathnbr	decktypeid	...	\
0	5.570193	3.5	66.0	...	
1	5.570193	1.0	66.0	...	
2	7.000000	2.0	66.0	...	
3	1.000000	4.0	66.0	...	
4	7.000000	1.0	66.0	...	
...	
31720	7.000000	1.0	66.0	...	
31721	7.000000	2.0	66.0	...	
31722	4.000000	3.0	66.0	...	
31723	7.000000	1.0	66.0	...	

31724	7.000000	1.0	66.0	...
-------	----------	-----	------	-----

	numberofstories	fireplaceflag	structuretaxvaluedollarcnt	\
0	1.450479	1	346458.0	
1	1.450479	1	66834.0	
2	1.450479	1	55396.0	
3	1.450479	1	559040.0	
4	1.450479	1	56233.0	
...	
31720	1.450479	1	65728.0	
31721	1.450479	1	40163.0	
31722	1.450479	1	248378.0	
31723	1.450479	1	17520.0	
31724	1.450479	1	66258.0	

	taxvaluedollarcnt	assessmentyear	landtaxvaluedollarcnt	taxamount	\
0	585529.0	2015	239071.0	10153.02	
1	210064.0	2015	143230.0	2172.88	
2	105954.0	2015	50558.0	1443.69	
3	1090127.0	2015	531087.0	13428.94	
4	70316.0	2015	14083.0	913.17	
...	
31720	307167.0	2015	241439.0	4038.70	
31721	50203.0	2015	10040.0	1263.39	
31722	331525.0	2015	83147.0	6461.79	
31723	39934.0	2015	22414.0	627.91	
31724	163037.0	2015	96779.0	2560.96	

	taxdelinquencyflag	taxdelinquencyyear	censustractandblock
0	NaN	13.314845	6.048996e+13
1	NaN	13.314845	6.059040e+13
2	NaN	13.314845	6.037140e+13
3	NaN	13.314845	6.037260e+13
4	NaN	13.314845	6.037570e+13
...
31720	NaN	13.314845	6.037550e+13
31721	Y	15.000000	6.037900e+13
31722	NaN	13.314845	6.037900e+13
31723	NaN	13.314845	6.037230e+13
31724	NaN	13.314845	6.037540e+13

[31725 rows x 58 columns]

```
[68]: # split and train
X = df.select_dtypes('float').values[:,1:]
y = df.select_dtypes('float').values[:,0]
```



```

cv = KFold(n_splits=5, random_state=1, shuffle=True) #default is 5 (redundant)
scores = cross_val_score(model, X_cv, y_cv, scoring='neg_mean_squared_error',
    ↪cv=cv, n_jobs=-1)
print('Mean Squared Error_CV:', np.mean(np.absolute(scores)))

```

Mean Squared Error_CV: 0.014292075903614952

```

[71]: # Run d2 for 100 times
import random

for i in range(100):
    seed = random.randint(0,99)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=seed)

    #     sc = StandardScaler()
    #     X_train = sc.fit_transform(X_train)
    #     X_test = sc.transform(X_test)

    regressor = RandomForestRegressor(n_estimators=20, random_state=seed)
    regressor.fit(X_train, y_train)
    y_prd = regressor.predict(X_test)
    print('random_seed: ',seed,'MSE:', metrics.mean_squared_error(y_test,
    ↪y_prd))

```

```

random_seed: 49 MSE: 0.010374547672181215
random_seed: 64 MSE: 0.010468170321508635
random_seed: 27 MSE: 0.010738558893648883
random_seed: 49 MSE: 0.010374547672181215
random_seed: 40 MSE: 0.010652433594255526
random_seed: 29 MSE: 0.010615661260027666
random_seed: 49 MSE: 0.010374547672181215
random_seed: 27 MSE: 0.010738558893648883
random_seed: 18 MSE: 0.010661646303790831
random_seed: 75 MSE: 0.010056237810833929
random_seed: 72 MSE: 0.010346488533182234
random_seed: 7 MSE: 0.010596493018438952
random_seed: 91 MSE: 0.010704934797804322
random_seed: 43 MSE: 0.010759998565866298
random_seed: 38 MSE: 0.010741879316798836
random_seed: 77 MSE: 0.010849221227234375
random_seed: 84 MSE: 0.010447766787825772
random_seed: 69 MSE: 0.010847690855126105
random_seed: 49 MSE: 0.010374547672181215
random_seed: 9 MSE: 0.010424234396700042
random_seed: 61 MSE: 0.010612658205591279
random_seed: 23 MSE: 0.010710105239384803

```

random_seed: 93 MSE: 0.010859594909032997
random_seed: 98 MSE: 0.010582599725900184
random_seed: 31 MSE: 0.010980465063084267
random_seed: 35 MSE: 0.010644306478161933
random_seed: 83 MSE: 0.010759286545648037
random_seed: 48 MSE: 0.010803725674328008
random_seed: 30 MSE: 0.01087953889017812
random_seed: 29 MSE: 0.010615661260027666
random_seed: 66 MSE: 0.010624729789045708
random_seed: 13 MSE: 0.010856756087537044
random_seed: 66 MSE: 0.010624729789045708
random_seed: 32 MSE: 0.01085643501810313
random_seed: 53 MSE: 0.010413282115624773
random_seed: 29 MSE: 0.010615661260027666
random_seed: 35 MSE: 0.010644306478161933
random_seed: 36 MSE: 0.011157417113462553
random_seed: 6 MSE: 0.010933839022101186
random_seed: 88 MSE: 0.010751002850993351
random_seed: 99 MSE: 0.010971654293104337
random_seed: 25 MSE: 0.010753152963904783
random_seed: 7 MSE: 0.010596493018438952
random_seed: 35 MSE: 0.010644306478161933
random_seed: 40 MSE: 0.010652433594255526
random_seed: 9 MSE: 0.010424234396700042
random_seed: 46 MSE: 0.010307968058643967
random_seed: 38 MSE: 0.010741879316798836
random_seed: 19 MSE: 0.010716984213912935
random_seed: 57 MSE: 0.010768774084516734
random_seed: 90 MSE: 0.010381302678966874
random_seed: 34 MSE: 0.011247187826096326
random_seed: 91 MSE: 0.010704934797804322
random_seed: 92 MSE: 0.010708301653129328
random_seed: 26 MSE: 0.01083831778438862
random_seed: 84 MSE: 0.010447766787825772
random_seed: 19 MSE: 0.010716984213912935
random_seed: 60 MSE: 0.010591942493656225
random_seed: 15 MSE: 0.010554539942102875
random_seed: 83 MSE: 0.010759286545648037
random_seed: 22 MSE: 0.010287365101144821
random_seed: 90 MSE: 0.010381302678966874
random_seed: 89 MSE: 0.010908196113060164
random_seed: 72 MSE: 0.010346488533182234
random_seed: 91 MSE: 0.010704934797804322
random_seed: 78 MSE: 0.011250073584658234
random_seed: 77 MSE: 0.010849221227234375
random_seed: 23 MSE: 0.010710105239384803
random_seed: 69 MSE: 0.010847690855126105
random_seed: 19 MSE: 0.010716984213912935

```
random_seed: 73 MSE: 0.010286551802982683
random_seed: 53 MSE: 0.010413282115624773
random_seed: 51 MSE: 0.010775947624611643
random_seed: 76 MSE: 0.01061688470814898
random_seed: 64 MSE: 0.010468170321508635
random_seed: 86 MSE: 0.011083256725539977
random_seed: 84 MSE: 0.010447766787825772
random_seed: 25 MSE: 0.010753152963904783
random_seed: 99 MSE: 0.010971654293104337
random_seed: 41 MSE: 0.010952345344469934
random_seed: 93 MSE: 0.010859594909032997
random_seed: 29 MSE: 0.010615661260027666
random_seed: 16 MSE: 0.010477144279381526
random_seed: 60 MSE: 0.010591942493656225
random_seed: 96 MSE: 0.010924333303800698
random_seed: 87 MSE: 0.010554557590616645
random_seed: 58 MSE: 0.010330998457447272
random_seed: 60 MSE: 0.010591942493656225
random_seed: 29 MSE: 0.010615661260027666
random_seed: 70 MSE: 0.010577889454047267
random_seed: 61 MSE: 0.010612658205591279
random_seed: 6 MSE: 0.010933839022101186
random_seed: 56 MSE: 0.010982175440280791
random_seed: 63 MSE: 0.010534524775824654
random_seed: 7 MSE: 0.010596493018438952
random_seed: 62 MSE: 0.010164430815652757
random_seed: 49 MSE: 0.010374547672181215
random_seed: 70 MSE: 0.010577889454047267
random_seed: 16 MSE: 0.010477144279381526
random_seed: 9 MSE: 0.010424234396700042
```

There is not a huge different between the MSE with different random seeds. Cross_validation would be good here because it would build K different models so we are able to make prediction on all the data. Confidence in the alogrithm would also be built with cross-validation as confidence in the algorithm is currently weak.

[]:

[]: